

# The future with AI (5)

15 feb., 18:30 – 20:00

Cluj-Napoca, Strada Teodor Mihali 62

> Let's talk about "prompt engineering"



**curs-mi.com**  
Machine Learning Course

**TECH 'N TRADE**



# Agenda

- Introduction
- Community news
- AI News
- Tech review
- (Break)
- Let's talk about “prompt engineering”
- (Networking)



# Community News

# Community news

- [GitHub repository](#) with all the codes and presentation
- [Discord server](#) (free to join) - <https://discord.gg/qd687uSW>



# GitHub



**curs-ml.com**  
Machine Learning Course

**TECH 'N' TRADE**

AI news



# AI News

- [DeepSeek Coder](#)
- [OpenLLMetry-JS](#)
- [MiQU-1 70b \(Mixtral Medium Model\)](#)
- Self Rewarding Language Models
- BioDrone (China)
- Qwen-VL (China)



# DeepSeek Coder

- [GitHub Repo](#)
- Copilot Alternative
- Demo



# OpenLLMetry-JS

- [GitHub Repo](#)



+



+





# MiQU-1 70b (Mistral Medium)

- [MiQU-1 70b \(Mistral Medium Model\)](#)

Rank ▲	🏆 Model ▲	🌟 Arena Elo ▲	🇮🇹 95% CI ▲	🗳 Votes ▲	Organization
1	<a href="#">GPT-4-0125-preview</a>	1253	+10/-11	3922	OpenAI
2	<a href="#">GPT-4-1106-preview</a>	1252	+5/-6	35385	OpenAI
3	<a href="#">Bard (Gemini Pro)</a>	1224	+9/-9	9081	Google
4	<a href="#">GPT-4-0314</a>	1190	+5/-6	18945	OpenAI
5	<a href="#">GPT-4-0613</a>	1162	+4/-5	29950	OpenAI
6	<a href="#">Mistral Medium</a>	1150	+6/-7	15447	Mistral
7	<a href="#">Claude-1</a>	1149	+6/-6	18189	Anthropic
8	<a href="#">Claude-2.0</a>	1132	+6/-7	12131	Anthropic



# Self-Rewarding Language Models

- RLHF without the RL (Facebook)
  - Classification with HF (DPO)
  - Classification with LLM Feedback (SRLM)

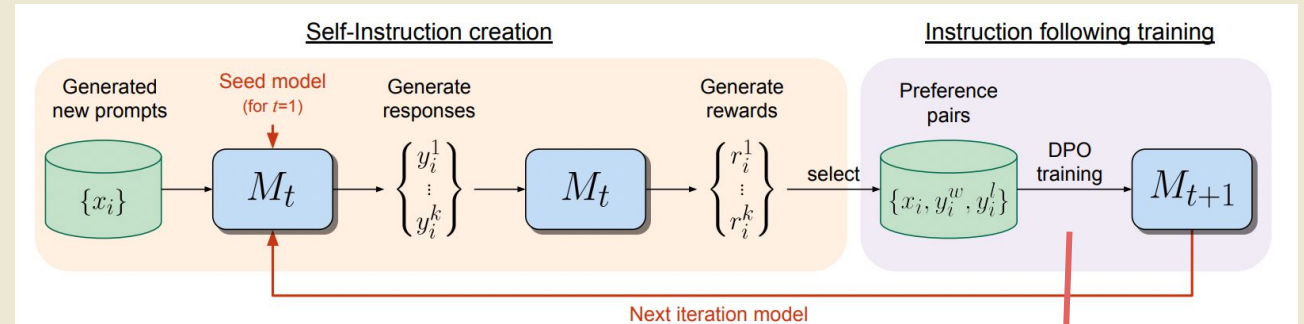


Figure 1: **Self-Rewarding Language Models.** Our self-alignment method consists of two steps: (i) *Self-Instruction creation*: newly created prompts are used to generate candidate responses from model  $M_t$ , which also predicts its own rewards via LLM-as-a-Judge prompting. (ii) *Instruction following training*: preference pairs are selected from the generated data, which are used for training via DPO, resulting in model  $M_{t+1}$ . This whole procedure can then be iterated resulting in both improved instruction following and reward modeling ability.





Figure 1: **DPO optimizes for human preferences while avoiding reinforcement learning.** Existing methods for fine-tuning language models with human feedback first fit a reward model to a dataset of prompts and human preferences over pairs of responses, and then use RL to find a policy that maximizes the learned reward. In contrast, DPO directly optimizes for the policy best satisfying the preferences with a simple classification objective, fitting an *implicit* reward model whose corresponding optimal policy can be extracted in closed form.



# BioDrone (China)

- [GitHub Repo](#)
- Drone Dataset (142 GB)





# Qwen-VL (China)

- [GitHub Repo](#)
- Visual Language Model (GPT4-V)
- [Demo](#)

Model	DocVQA	ChartQA	AI2D	TextVQA	MMMU	MathVista	MM-Bench-CN
Other Best Open-source LVLM	81.6% (CogAgent)	68.4% (CogAgent)	73.7% (Fuyu-Medium)	76.1% (CogAgent)	45.9% (Yi-VL-34B)	36.7% (SPHINX-V2)	72.4% (InternLM-XComposer-VL)
Gemini Pro	88.1%	74.1%	73.9%	74.6%	47.9%	45.2%	74.3%
Gemini Ultra	90.9%	80.8% <sup>1</sup>	79.5% <sup>1</sup>	82.3% <sup>1</sup>	59.4% <sub>1</sub>	53.0% <sup>1</sup>	-
GPT-4V	88.4%	78.5%	78.2%	78.0%	56.8%	49.9%	73.9%
Qwen-VL-Plus	91.4%	78.1%	75.9%	78.9%	45.2%	43.3%	68.0%
Qwen-VL-Max	93.1% <sup>1</sup>	79.8% <sup>2</sup>	79.3% <sup>2</sup>	79.5% <sup>2</sup>	51.4% <sub>3</sub>	51.0% <sup>2</sup>	75.1% <sup>1</sup>



Tech news

# Ray-Ban Meta

- Spy camera glasses
  - Photos
  - Videos
- Live-streaming (on Facebook)
- Speakers / Microphone
- Llama 2 (AI) Assistant
- (not sold in EU...)



(Break)

Let's talk about  
“prompt engineering”



# Everyone does LLM these days..



# Quick recap of an LLM


## Next-Token Prediction in Math

$$\mathcal{U} = \{u_1, u_2, \dots, u_N\}$$
$$\mathcal{L}(\mathcal{U}) = \sum_{i=1}^N \log (\mathbb{P}(u_i | u_{i-k}, \dots, u_{i-1}, \Theta))$$

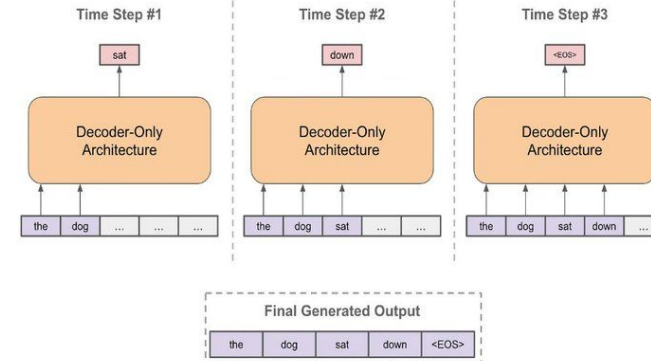
Language model loss over the full text corpus

Conditional probability of i-th token given k preceding tokens and model parameters  $\Theta$

## Pre-Training with Next-Token Prediction



## Autoregressive Decoding



# What is the “Prompt”

- It's the input text.
- But not all input texts are equal.

“All ~~models~~ prompts are wrong, but some are useful.”  
- ~~George Box~~ Cristian Lungu

“THE BAD ARTISTS  
IMITATE, THE GREAT  
ARTISTS STEAL.”

~~PABLO PICASSO~~  
BANKSY



# (correct) Prompt format

```
<s> [INST] QUERY_1 [/INST] ANSWER_1</s> [INST] QUERY_2 [/INST]  
ANSWER_2</s>...
```

- [Specific to each model](#)



# Use-case: Classification (ChatGPT)

```
{job_posting}
```

```
-----
```

```
Is this job (A) a job fit for a recent graduate, or  
(B) a job requiring more professional experience?"""
```



# Use-case: Classification

```
system = "You are Frederick, an AI expert in career advice. You  
are tasked with sorting through jobs by analysing their content  
and deciding whether they would be a good fit for a recent  
graduate or not."
```

```
user 1 = """A job is fit for a graduate if it's a  
junior-level position that does not require extensive prior  
professional experience.  
When analysing the experience required, take into account that  
requiring internships is still fit for a graduate. I will  
give you a job posting and you will analyse it, step-by-step,  
to know whether or not it describes a position fit for a  
graduate. Got it?"""
```

```
assistant 1 = "Yes, I understand. I am Frederick, and I will  
analyse your job posting."
```

```
user 2 = """Great! Let's begin then :)  
For the given job:  
{job_posting}
```

```
-----
```

```
Is this job (A) a job fit for a recent graduate, or  
(B) a job requiring more professional experience.  
Answer: Let's think step by step to reach the right  
conclusion"""
```



# Use-case: Classification

```
system = "You are Frederick, an AI expert in career advice. You  
are tasked with sorting through jobs by analysing their content  
and deciding whether they would be a good fit for a recent  
graduate or not."
```

```
user 1 = """A job is fit for a graduate if it's a  
junior-level position that does not require extensive prior  
professional experience.  
When analysing the experience required, take into account that  
requiring internships is still fit for a graduate.  
I will give you a job posting and you will analyse it,  
step-by-step, to know whether or not it describes a position fit  
for a graduate. Got it?"""
```

```
assistant 1 = "Yes, I understand. I am Frederick, and I will  
analyse your job posting."
```

```
user 2 = """Great! Let's begin then :)  
For the given job:  
{job_posting}  
-----  
Is this job (A) a job fit for a recent graduate, or  
(B) a job requiring more professional experience.  
Answer: Let's think step by step to reach the right  
conclusion"""
```



# Use-case: Classification

- The baseline prompt without any decorations
- Persona definition** – Define the role / character this AI will play
- Task definition** – Define the task this AI will solve
- Split prompt** – Use both the *system* and *user* sections. *System* contains the imperative "who you are".
- Zero-shot CoT (Chain of thought)** – Prime the assistant as it is
  - Name** – Give the LLM a name
  - Reiterate** – Repeat certain main, important instructions.
  - Positive Feedback** – Reassure that the above instructions are correct and to be followed.
  - Edge cases** – ...
  - Be right!** – Ask the model to find the correct solution, not *any* solution
- Mocked Exchange** – Break down on the prompt into a mocked conversation that the agent should be bound to





# Use-case: Classification

```
system = """You are Frederick, an AI expert in career advice.
You are tasked with sorting through jobs by
analysing their content and deciding whether
they would be a good fit for a recent graduate or not."""
user 1 = """A job is fit for a graduate if it's a junior-level
position that does not require extensive prior
professional experience.
When analysing the experience required, take
into account that requiring internships is still
fit for a graduate.
I will give you a job posting and you will analyse
it, step-by-step, to know whether or not it
describes a position fit for a graduate.
Got it?"""
assistant 1 = """Yes, I understand. I am Frederick, and I will analyse
your job posting."""
user 2 = """Great! Let's begin then :)

For the given job:
{job_posting}
-----

Is this job (A) a job fit for a recent graduate,
or (B) a job requiring more professional experience.
Answer: Let's think step by step
to reach the right conclusion"""
```

Persona definition

Task definition

Split prompt

Zero-shot CoT (Chain of thought)

Name

Reiterate

Positive Feedback

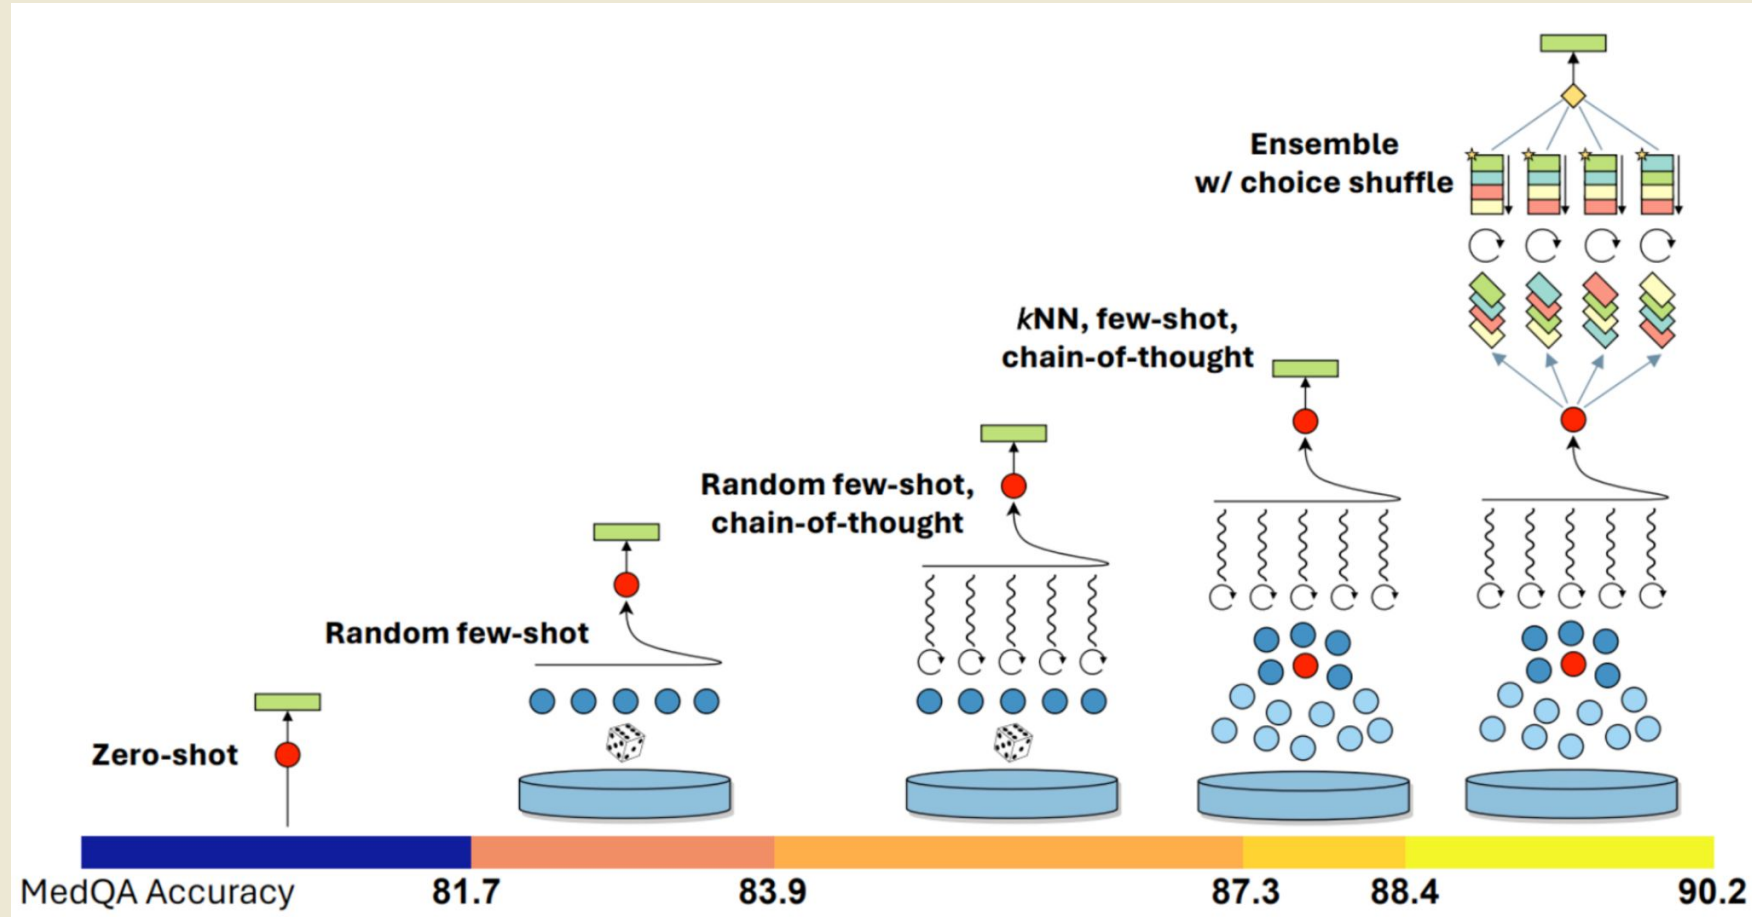
Edge cases

Be right!

Mocked Exchange



# Use-case: MedQA



Why do we need to this?



# Because of instruction tuning


```
### Instruction:  
{instruction}
```

```
### Input:  
{input}
```

```
### Response:
```



## Because of instruction tuning (Alpaca)



# Reverse engineering the language patterns of the instructors





# Prompt Engineering is NOT Engineering







# Use-Case: LLM Guardrails - Instructor

```
import instructor
from openai import OpenAI
from pydantic import BaseModel

# Enables `response_model`
client = instructor.patch(OpenAI())

class UserDetail(BaseModel):
    name: str
    age: int

user = client.chat.completions.create(
    model="gpt-3.5-turbo",
    response_model=UserDetail,
    messages=[
        {"role": "user", "content": "Extract Jason is 25 years old"},
    ],
)

assert isinstance(user, UserDetail)
assert user.name == "Jason"
assert user.age == 25
```



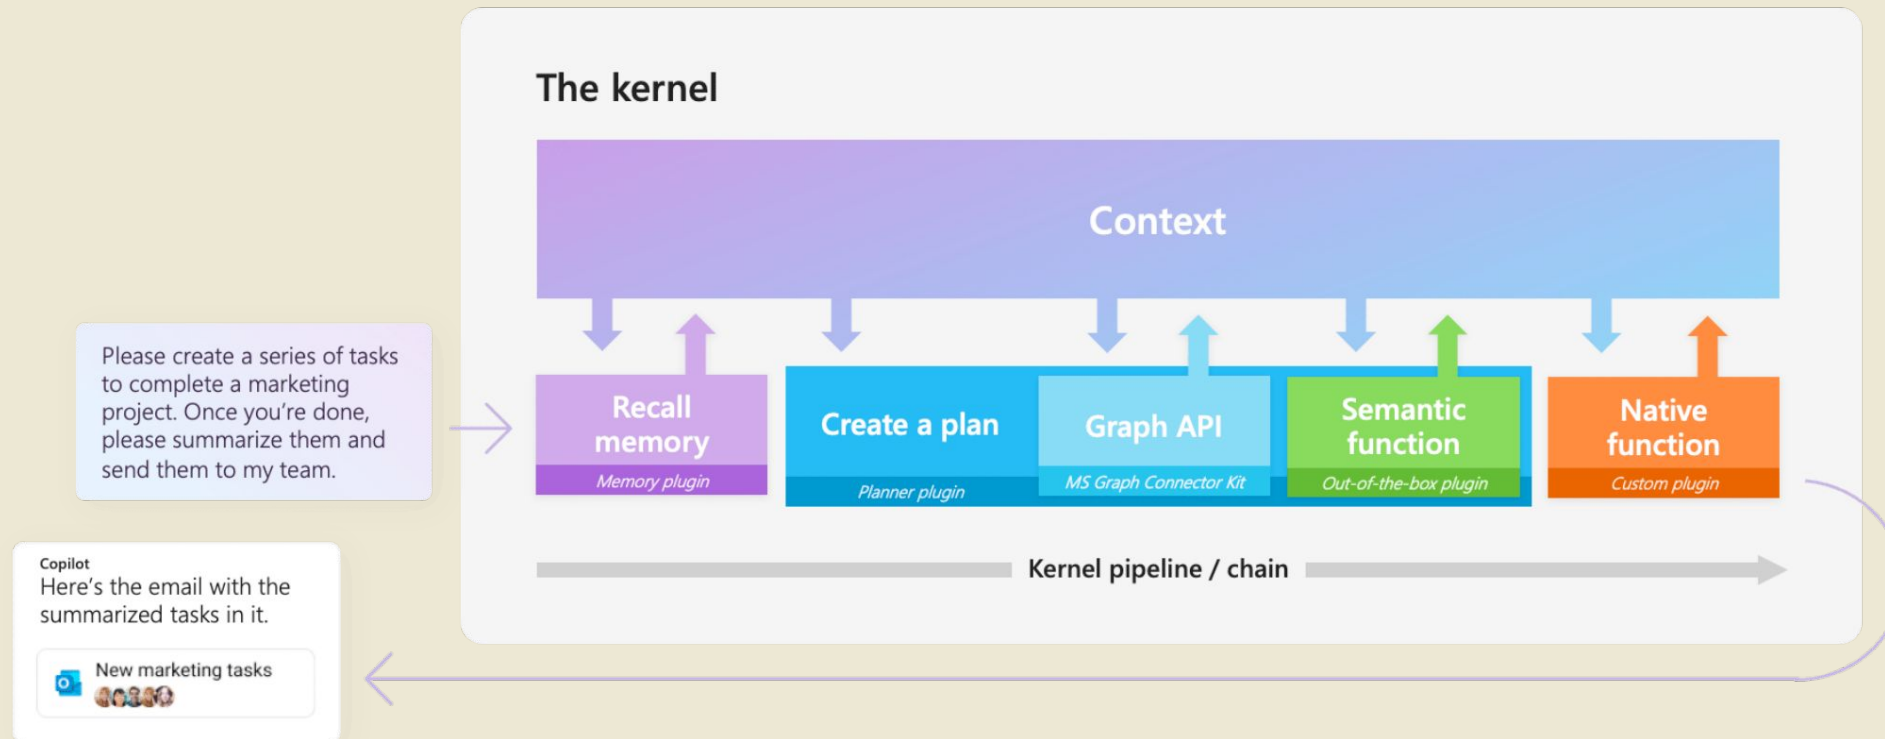
# Use-Case: LLM Guardrails - Instructor

```
if mode == Mode.TOOLS:
    kwargs["messages"].append(
        {
            "role": "tool",
            "tool_call_id": response.choices[0]
                .message.tool_calls[0]
                .id,
            "name": response.choices[0]
                .message.tool_calls[0]
                .function.name,
            "content": f"Recall the function correctly, fix the errors and exceptions found\n{e}",
        }
    )
else:
    kwargs["messages"].append(
        {
            "role": "user",
            "content": f"Recall the function correctly, fix the errors and exceptions found\n{e}",
        }
    )

if mode == Mode.MD_JSON:
    kwargs["messages"].append(
        {
            "role": "assistant",
            "content": "```json",
        },
    )
```



# Use-Case: Semantic Kernel



## Use-Case: Semantic Kernel

```
{{#*inline "RetryLogic"}}
{{~#if lastError}}
{{#message role="system"}}## Previous attempt
This previous plan failed to achieve the goal:
```handlebars
{{lastPlan}}
```

The error was:
```
{{lastError}}
```

Try again to achieve the goal while fixing the error.
{{/message}}
{{/if}}
{{/inline}}
```



# Other “esoteric” techniques

- Sparse Language Representation
  - This is a technique for "priming" the LLM in the context prompt (system prompt) with specific words that would allow the attention mechanism for promoting specific **paths** from the llm memory.
  - This is basically a way to condense into the context window (in a token efficient way) information about what the original system prompt is.
- Dynamic Few Shots
- ...



# Conclusions

- Every model has a different “best prompt”
- Prompts are brittle
- All “best practices” are just hacks that try to mimic the training dataset
- This is not engineering



# Thank you!



**curs-mi.com**  
Machine Learning Course

**TECH 'N TRADE**